

# assignment0

January 14, 2026

## 1 CS260R Reinforcement Learning Assignment 0: Jupyter Notebook usage and assignment submission workflow

---

*CS260R: Reinforcement Learning. Department of Computer Science at University of California, Los Angeles. Course Instructor: Professor Bolei ZHOU. Assignment author: Zhenghao PENG.*

You are asked to finish four tasks:

1. Fill in your name and University ID in the next cell.
2. Install pytorch and finish the [Kindergarten Pytorch](#) section.
3. Run all cells and export this notebook [as a PDF file](#).
4. Compress this folder `assignment0` [as a ZIP file](#) and submit [the PDF file](#) and [the ZIP file separately](#) as two files in BruinLearn.

```
[1]: # TODO: Fill your name and UID here
my_name = "Anthony Zhao"
my_student_id = "906173995"
```

```
[2]: # Run this cell without modification

text = "Oh, I finished this assignment! I am {} ({}).format(my_name, my_student_id)
print(text)
with open("{} .txt".format(text), "w") as f:
    f.write(text)
```

Oh, I finished this assignment! I am Anthony Zhao (906173995)

### 1.1 Kindergarten Pytorch

1. Please install pytorch in your virtual environment following the instruction: <https://pytorch.org/get-started/locally/>.

```
pip install torch torchvision
```

2. If you are not familiar with Pytorch, please go through [the tutorial in official website](#) until you can understand the [quick start tutorial](#).
3. The following code is copied from the [quick start tutorial](#), please solve all TODOs and print the result in the cells before generating the PDF file.

```
[3]: # You can also run this cell in notebook:  
# !pip install torch torchvision
```

### 1.1.1 Prepare data

```
[4]: import torch  
from torch import nn  
from torch.utils.data import DataLoader  
from torchvision import datasets  
from torchvision.transforms import ToTensor  
  
# Download training data from open datasets.  
training_data = datasets.FashionMNIST(  
    root="data",  
    train=True,  
    download=True,  
    transform=ToTensor(),  
)  
  
# Download test data from open datasets.  
test_data = datasets.FashionMNIST(  
    root="data",  
    train=False,  
    download=True,  
    transform=ToTensor(),  
)  
  
batch_size = 64  
  
# Create data loaders.  
train_dataloader = DataLoader(training_data, batch_size=batch_size)  
test_dataloader = DataLoader(test_data, batch_size=batch_size)
```

### 1.1.2 Define model

```
[5]: # Get cpu, gpu or mps device for training.  
device = (  
    "cuda"  
    if torch.cuda.is_available()  
    else "mps"  
    if torch.backends.mps.is_available()  
    else "cpu"  
)  
print(f"Using {device} device")  
  
# Define model
```

```

class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()

        # TODO: Define the self.linear_relu_stack by uncommenting next few lines
        # and understand what they mean
        # Neural net with 2 hidden layers (512 neurons each) and 1 output layer
        # (10 neurons) and 1 input layer (28*28)
        # there are 10 classes of articles of clothing and each image is
        # transformed into a 28x28 pixel image
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28 * 28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10)
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

model = NeuralNetwork().to(device)
print(model)

```

Using mps device

```

NeuralNetwork(
    (flatten): Flatten(start_dim=1, end_dim=-1)
    (linear_relu_stack): Sequential(
        (0): Linear(in_features=784, out_features=512, bias=True)
        (1): ReLU()
        (2): Linear(in_features=512, out_features=512, bias=True)
        (3): ReLU()
        (4): Linear(in_features=512, out_features=10, bias=True)
    )
)

```

### 1.1.3 Define training and test pipelines

```
[6]: loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=1e-3)

def train(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)
```

```

model.train()
for batch, (X, y) in enumerate(dataloader):
    X, y = X.to(device), y.to(device)

    # Compute prediction error
    pred = model(X)
    loss = loss_fn(pred, y)

    # Backpropagation

    # TODO: Uncomment next three lines and understand what they mean
    loss.backward()
    optimizer.step()
    optimizer.zero_grad()

    if batch % 100 == 0:
        loss, current = loss.item(), (batch + 1) * len(X)
        print(f"loss: {loss:.7f} [{current:.5d}/{size:.5d}]")

def test(dataloader, model, loss_fn):
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    model.eval()
    test_loss, correct = 0, 0
    with torch.no_grad():
        for X, y in dataloader:
            X, y = X.to(device), y.to(device)
            pred = model(X)
            test_loss += loss_fn(pred, y).item()

            # TODO: Uncomment next line and understand what it means
            # "how many predictions are correct in the batch"
            # pred.argmax(1) is the predicted class of the batch
            # y is the actual class of the batch
            # so (pred.argmax(1) == y) is a boolean array of the correct
            # predictions of size batch_size
            # .type(torch.float) converts the boolean array to a float array of
            # 0s and 1s
            # .sum() sums the float array to get the number of correct
            # predictions
            # .item() gets the number of correct predictions as a scalar
            correct += (pred.argmax(1) == y).type(torch.float).sum().item()

    test_loss /= num_batches
    correct /= size
    print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss:\
{test_loss:>8f} \n")

```

#### 1.1.4 Run the training and test pipelines

```
[7]: epochs = 5
for t in range(epochs):
    print(f"Epoch {t+1}\n-----")
    train(train_dataloader, model, loss_fn, optimizer)
    test(test_dataloader, model, loss_fn)
print("Done!")
```

Epoch 1

```
-----
loss: 2.298511  [  64/60000]
loss: 2.290683  [ 6464/60000]
loss: 2.267752  [12864/60000]
loss: 2.267349  [19264/60000]
loss: 2.248700  [25664/60000]
loss: 2.217807  [32064/60000]
loss: 2.228979  [38464/60000]
loss: 2.195093  [44864/60000]
loss: 2.187416  [51264/60000]
loss: 2.164214  [57664/60000]
Test Error:
Accuracy: 45.4%, Avg loss: 2.150111
```

Epoch 2

```
-----
loss: 2.157988  [  64/60000]
loss: 2.147011  [ 6464/60000]
loss: 2.080500  [12864/60000]
loss: 2.104212  [19264/60000]
loss: 2.049226  [25664/60000]
loss: 1.986271  [32064/60000]
loss: 2.019812  [38464/60000]
loss: 1.935037  [44864/60000]
loss: 1.937029  [51264/60000]
loss: 1.878115  [57664/60000]
Test Error:
Accuracy: 55.6%, Avg loss: 1.861714
```

Epoch 3

```
-----
loss: 1.893943  [  64/60000]
loss: 1.862058  [ 6464/60000]
loss: 1.733222  [12864/60000]
loss: 1.782529  [19264/60000]
loss: 1.678723  [25664/60000]
loss: 1.628315  [32064/60000]
loss: 1.655413  [38464/60000]
```

```
loss: 1.559976 [44864/60000]
loss: 1.583152 [51264/60000]
loss: 1.491168 [57664/60000]
Test Error:
Accuracy: 62.5%, Avg loss: 1.499361
```

Epoch 4

```
-----  
loss: 1.563573 [ 64/60000]
loss: 1.535503 [ 6464/60000]
loss: 1.379463 [12864/60000]
loss: 1.451912 [19264/60000]
loss: 1.345607 [25664/60000]
loss: 1.336782 [32064/60000]
loss: 1.348679 [38464/60000]
loss: 1.283909 [44864/60000]
loss: 1.315861 [51264/60000]
loss: 1.225344 [57664/60000]
Test Error:
Accuracy: 64.3%, Avg loss: 1.244303
```

Epoch 5

```
-----  
loss: 1.317363 [ 64/60000]
loss: 1.307872 [ 6464/60000]
loss: 1.137531 [12864/60000]
loss: 1.237312 [19264/60000]
loss: 1.123834 [25664/60000]
loss: 1.142267 [32064/60000]
loss: 1.157967 [38464/60000]
loss: 1.108679 [44864/60000]
loss: 1.147247 [51264/60000]
loss: 1.066365 [57664/60000]
Test Error:
Accuracy: 65.3%, Avg loss: 1.081602
```

Done!

### 1.1.5 Save model

```
[8]: torch.save(model.state_dict(), "model.pth")
print("Saved PyTorch Model State to model.pth")
```

Saved PyTorch Model State to model.pth

### 1.1.6 Load model and run the inference

```
[9]: model = NeuralNetwork().to(device)
model.load_state_dict(torch.load("model.pth"))
```

```
[9]: <All keys matched successfully>
```

```
[10]: classes = [
        "T-shirt/top",
        "Trouser",
        "Pullover",
        "Dress",
        "Coat",
        "Sandal",
        "Shirt",
        "Sneaker",
        "Bag",
        "Ankle boot",
    ]

model.eval()
x, y = test_data[0][0], test_data[0][1]
with torch.no_grad():
    x = x.to(device)
    pred = model(x)
predicted, actual = classes[pred[0].argmax(0)], classes[y]
print(f'Predicted: "{predicted}", Actual: "{actual}"')
```

```
Predicted: "Ankle boot", Actual: "Ankle boot"
```